

Multivariate Regressions

Michele Coscia

First Year Project #1

February 17th, 2022

Lecture Plan

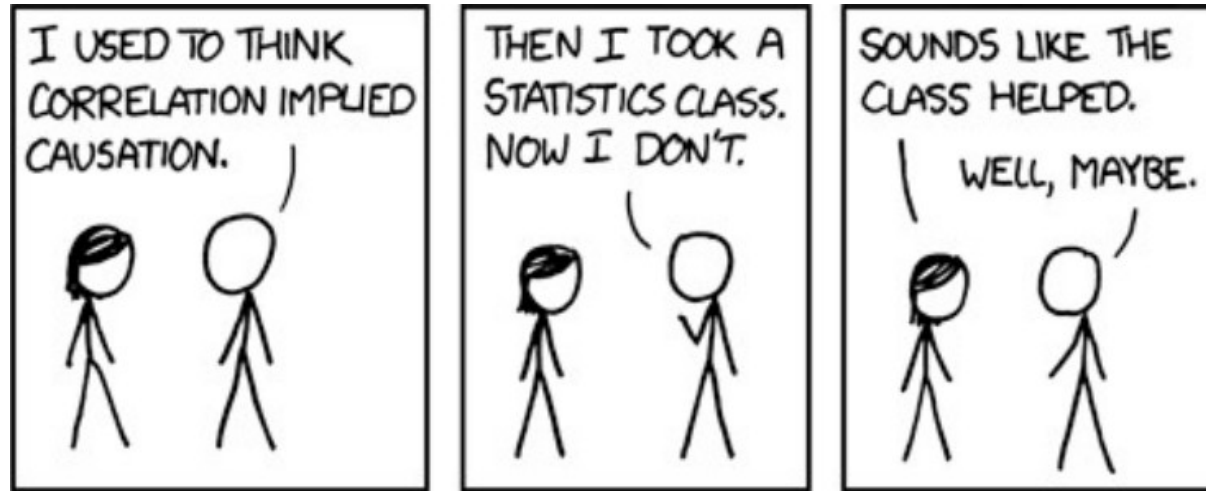
- 1) (February 8th) Intro
- 2) (February 10th) Geospatial Basics
- 3) (February 15th) Estimating Associations
- 4) (Today) Multivariate Regression**
- 5) (February 22nd) Interventions
- 6) (February 24th) Project Run Through
- 7) (March 1st) Q&A – Open Supervision
- 8) (March 3rd) Q&A – Open Supervision

Outline

- Multivariate Regression Basics
- Fixed Effects
- Clustered Standard Errors

Multivariate Regression Basics

Always Remember



Correlation \neq Causation

- Because X & Y could be both dependent on Z
- Pearson cannot handle this
- Multivariate regression can!
- You can control for Z
- \rightarrow Investigate the effect of X keeping Z constant

Multivariate regression is still not enough!

- Linear, just like Pearson
- Finding all Zs is hard
- Some controls are bad!
- It won't tell you if $X \rightarrow Y$ or $Y \rightarrow X$
- Causality is hard...

Linearity

- Running a regression comes with assumptions
- Non-linear data normally break these assumptions
- Hard way: check the assumption list and satisfy them
- For now: log-transform

Let's Get Started!

The screenshot displays a JupyterLab environment with a file browser on the left and a notebook editor on the right. The notebook contains two code cells. The first cell imports necessary libraries: `json`, `numpy`, `pandas`, `statsmodels.api`, and `scipy.stats`. The second cell performs data cleaning and merging, reading CSV files for COVID-19 data, metadata, and weather, then merging them into a single DataFrame `df`.

Below the code, the output of the second cell is shown as a table with 12 columns: `date`, `hospitalized_addition`, `population`, `cases_pc`, `iso3166-2`, `RelativeHumiditySurface`, `SolarRadiation`, `Surfacepressure`, `TemperatureAboveGround`, `Totalprecipitation`, `UVindex`, and `WindSpeed`.

	date	hospitalized_addition	population	cases_pc	iso3166-2	RelativeHumiditySurface	SolarRadiation	Surfacepressure	TemperatureAboveGround	Totalprecipitation	UVindex	WindSpeed
0	2020-03-01	1	1846023	5.417051e-07	DK-84	79.371362	3.383109e+06	2.370635e+06	5.064128	0.000764	2.595843	6.668466
1	2020-03-02	0	1846023	0.000000e+00	DK-84	86.574612	3.303007e+06	2.380293e+06	4.470362	0.001416	4.286374	2.475038
2	2020-03-03	1	1846023	5.417051e-07	DK-84	93.285949	9.690623e+04	2.395165e+06	3.884757	0.002084	1.676674	2.345198
3	2020-03-04	0	1846023	0.000000e+00	DK-84	86.105840	3.227602e+06	2.407377e+06	4.677848	0.000926	4.771363	4.631544
4	2020-03-05	1	1846023	5.417051e-07	DK-84	86.688654	2.998848e+06	2.403363e+06	3.949029	0.000420	4.919169	2.801289
...
1755	2021-02-11	1	589936	1.695099e-06	DK-81	73.558470	3.624393e+06	2.475768e+06	-6.216205	0.000383	1.495042	4.113037
1756	2021-02-12	1	589936	1.695099e-06	DK-81	74.618363	4.379149e+06	2.491939e+06	-6.035219	0.000006	1.992372	1.915713
1757	2021-02-13	1	589936	1.695099e-06	DK-81	76.532522	4.910543e+06	2.494230e+06	-4.408170	0.000000	2.279176	1.357024
1758	2021-02-14	1	589936	1.695099e-06	DK-81	74.459283	4.752374e+06	2.484782e+06	-3.379998	0.000000	2.772693	2.861502
1759	2021-02-15	2	589936	3.390198e-06	DK-81	76.951013	3.486211e+04	2.452897e+06	-0.972067	0.002582	0.015256	5.553239

The bottom status bar indicates the interface is in 'Simple' mode, running Python 3 (ipykernel) in an idle state. The current file is `pr01_e04.ipynb`, and the cursor is at line 2, column 52.

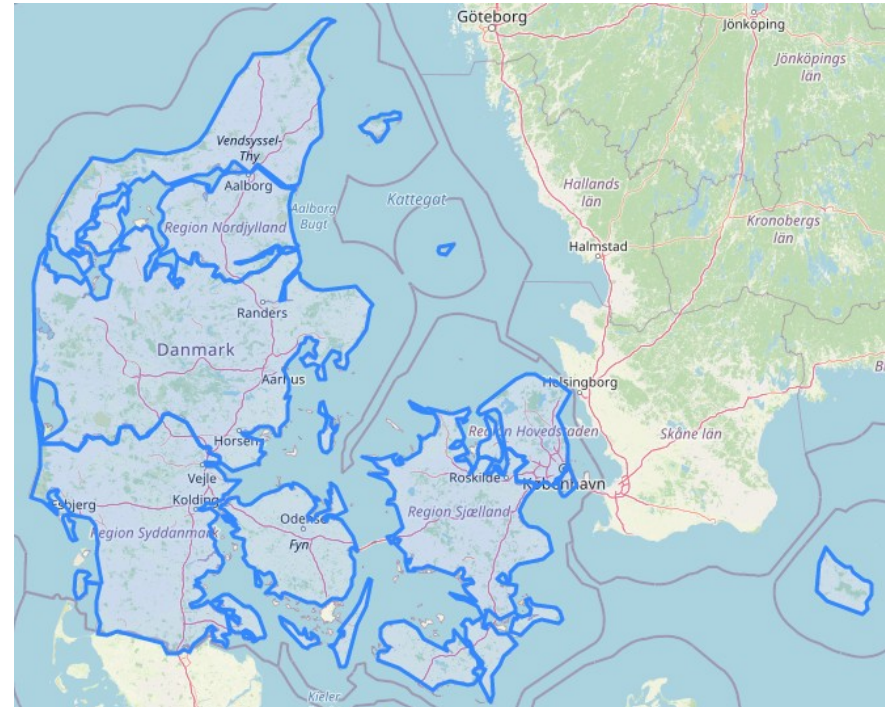
Fixed Effects

Problem

- Sometimes you **know** something affected your outcome
- You just don't have any measure for it
- In our case: different local governments work differently


Fixed Effects

- You know your observations belong to specific groups
 - In our case, Danish regions
- The avg of each group is fixed
- Everything that group does differently from the other groups is captured here




Why would we do this?

- Corr $\sim 0.8!!$
- Best fit
- Something Fishy...



Why would we do this?

- Groups!
- Related with X
- The true relationship is actually negative!



How to do it practically

- In R, it's automatic
 - Just pass a categorical variable to your regression function
- In general, you can add a “dummy variable”
 - One variable per group
 - 1 if observation belongs to the group, 0 otherwise
 - You need to omit one group (the reference)

Interpretation

- Coefficient tells you the effect of being part of the group
 - Specifically: the difference between your group and the reference one
- If group membership is important for your question, you can interpret it
 - But careful, because you're absorbing everything!
- Most often, it's just a control → Ignore

Let's Fix Some Effects!

File Edit View Run Kernel Tabs Settings Help

pr01_e04.ipynb | leftovers.ipynb

Python 3 (ipykernel)

strong multicollinearity or other numerical problems.

```
[34]: # Here we add a "dummy" variable: a region fixed effect, identify which rows belong
# to which region. This dummy variable absorbs every possible omitted variable that
# distinguishes a region from all other regions.
regions = ["const",]

for region in set(df["iso3166-2"]):
    if region != "DK-81":
        df[region] = (df["iso3166-2"] == region).astype(int)
        regions.append(region)
    Xs.append(region)

df
```

```
[34]:
```

	const	date	hospitalized_addition	population	cases_pc	iso3166-2	RelativeHumiditySurface	SolarRadiation	Surfacepressure	TemperatureAboveGround	Totalprecipitation	UVIndex	WindSpeed	DK-83	DK-84	DK-85
0	1.0	2020-03-01	1	1846023	5.417051e-07	DK-84	79.371362	3.383109e+06	2.370635e+06	5.064128	0.000764	2.595843	6.668466	0	1	
1	1.0	2020-03-02	0	1846023	0.000000e+00	DK-84	86.574612	3.303007e+06	2.380293e+06	4.470362	0.001416	4.286374	2.475038	0	1	
2	1.0	2020-03-03	1	1846023	5.417051e-07	DK-84	93.285949	9.690623e+04	2.395165e+06	3.884757	0.002084	1.676674	2.345198	0	1	
3	1.0	2020-03-04	0	1846023	0.000000e+00	DK-84	86.105840	3.227602e+06	2.407377e+06	4.677848	0.000926	4.771363	4.631544	0	1	
4	1.0	2020-03-05	1	1846023	5.417051e-07	DK-84	86.688654	2.998848e+06	2.403363e+06	3.949029	0.000420	4.919169	2.801289	0	1	
...
1755	1.0	2021-02-11	1	589936	1.695099e-06	DK-81	73.558470	3.624393e+06	2.475768e+06	-6.216205	0.000383	1.495042	4.113037	0	0	
1756	1.0	2021-02-12	1	589936	1.695099e-06	DK-81	74.618363	4.379149e+06	2.491939e+06	-6.035219	0.000006	1.992372	1.915713	0	0	
1757	1.0	2021-02-13	1	589936	1.695099e-06	DK-81	76.532522	4.910543e+06	2.494230e+06	-4.408170	0.000000	2.279176	1.357024	0	0	
1758	1.0	2021-02-14	1	589936	1.695099e-06	DK-81	74.459283	4.752374e+06	2.484782e+06	-3.379998	0.000000	2.772693	2.861502	0	0	
1759	1.0	2021-02-15	2	589936	3.390198e-06	DK-81	76.951013	3.486211e+04	2.452897e+06	-0.972067	0.002582	0.015256	5.553239	0	0	

1760 rows x 17 columns

```
[30]: est = sm.OLS(np.log(df["cases_pc"] + 1), df[regions], hasconst = True).fit()
# Let's first see how regions did overall. No real differences, except maybe
# Brussels (BRU) doing poorly and East Flanders (VOV) doing well.
print(est.summary())
```


OLS Regression Results

```
=====
Dep. Variable:      cases_pc      R-squared:            0.097
Model:              OLS          Adj. R-squared:         0.095
Method:             Least Squares      F-statistic:         47.23
Date:               Wed, 02 Feb 2022    Prob (F-statistic):    9.40e-38
Time:               15:28:57           Log-Likelihood:       18306.
No. Observations:   1760             AIC:                 -3.660e+04
=====
```

Simple 0 4 Python 3 (ipykernel) | Idle Mode: Command Ln 1, Col 14 pr01_e04.ipynb


Clustered Standard Errors

What's Wrong Here?



DARN, NOT SIGNIFICANT.

WE NEED MORE DATA.
HAVE THEM EACH TRY
YELLING INTO THE MIC
A FEW MORE TIMES.




PERFECT!

ARE YOU *SURE*
WE'RE DOING
SLOPE HYPOTHESIS
TESTING RIGHT?




You're lying to your model!
You're saying you have 12 **independent** observations
Not true! You have 3

What's a standard error?




- How sure you are about your estimation

What's a standard error?




- How sure you are about your estimation
- More obs = more confidence

What's a standard error?



- How sure you are about your estimation
- More obs = more confidence
- But not if they're part of the same group!

Clustered Standard Errors



- A true new obs would be independent from the previous ones
- Doing CSEs can take this into account

Let's Cluster Some Errors!

The screenshot displays the JupyterLab environment. On the left, a sidebar contains a file browser with a search bar and a list of files/notebooks under the path '/notebooks/'. The files include 'leftovers.ip...', 'pr01_e01...', 'pr01_e02...', 'pr01_e03...', and 'pr01_e04...', each with its last modified time.

The main area shows a code editor with three tabs: 'Launcher', 'pr01_e04.ipynb', and 'leftovers.ipynb'. The active tab is 'pr01_e04.ipynb', which contains the following Python code:

```
[12]: est = sm.OLS(np.log(df["cases_pc"] + 1), df[Xs], hasconst = True).fit(cov_type = "cluster", cov_kws = {"groups": df["iso3166-2"]}, use_t = True)
print(est.summary())
```

The output of the code execution is displayed below the code cell. It includes the OLS Regression Results summary, which provides statistical measures such as R-squared, F-statistic, AIC, BIC, and standard errors for various predictors. The results are presented in a structured format with columns for coefficients, standard errors, t-values, p-values, and confidence intervals.

Q&A